

Gaussian processes

Alejandro Veloz

ORGANIZED BY:



Funded by
the European Union


COLLABORATORS:




Email: alejandro.veloz@uv.cl

Slides and Labs: <https://aavelozb.github.io/gp>


Motivation: non-linear regression




Motivation: non-linear regression




Motivation: non-linear regression



Motivation: non-linear regression



Motivation: non-linear regression



Outline

Making predictions with Gaussians

Gaussian processes for regression

Gaussian processes for classification


Gaussian processes and neural networks

Applications

Gaussian distribution

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$


$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Gaussian distribution

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$


$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Gaussian distribution

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$


$$\Sigma = \begin{bmatrix} 1 & 0.6 \\ 0.6 & 1 \end{bmatrix}$$



Gaussian distribution

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$


$$\Sigma = \begin{bmatrix} 1 & 0.4 \\ 0.4 & 1 \end{bmatrix}$$



Gaussian distribution

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$


$$\Sigma = \begin{bmatrix} 1 & 0.1 \\ 0.1 & 1 \end{bmatrix}$$



Gaussian distribution



$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$

$$\Sigma = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.8 & 0.6 & 0.4 \\ 0.9 & 1 & 0.9 & 0.8 & 0.6 \\ 0.8 & 0.9 & 1 & 0.9 & 0.8 \\ 0.6 & 0.8 & 0.9 & 1 & 0.9 \\ 0.4 & 0.6 & 0.8 & 0.9 & 1 \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & 0 & \dots \\ 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & \dots \\ 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & \dots \\ 0.3 & 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & \dots \\ \vdots & \ddots \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & 0 & \dots \\ 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & \dots \\ 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & \dots \\ 0.3 & 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & \dots \\ \vdots & \ddots \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & 0 & \dots \\ 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & \dots \\ 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & \dots \\ 0.3 & 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & \dots \\ \vdots & \ddots \end{bmatrix}$$





New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & 0 & \dots \\ 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & \dots \\ 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & \dots \\ 0.3 & 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & \dots \\ \vdots & \ddots \end{bmatrix}$$






New visualization

$$\Sigma = \begin{bmatrix} 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & 0 & \dots \\ 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & 0 & \dots \\ 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & 0 & \dots \\ 0.3 & 0.6 & 0.9 & 1 & 0.9 & 0.6 & 0.3 & 0.1 & \dots \\ \vdots & \ddots \end{bmatrix}$$



New visualization




$\Sigma =$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning

$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$




$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$



Marginalization and conditioning



$$p(\mathbf{y} \mid \Sigma) \propto \exp\left(-\frac{1}{2}\mathbf{y}^\top \Sigma^{-1} \mathbf{y}\right)$$

$$\Sigma = \begin{bmatrix} 1 & 0.7 \\ 0.7 & 1 \end{bmatrix}$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$\Sigma =$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$\Sigma =$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$$\Sigma =$$





Marginalization and conditioning

$\Sigma =$





Marginalization and conditioning

$\Sigma =$





Marginalization and conditioning



$$\Sigma =$$





Regression using Gaussians



Regression using Gaussians





Regression using Gaussians



Regression using Gaussians

$$\Sigma =$$



Gaussian process

For any finite set of points, this process defines a joint Gaussian:

$$p(\mathbf{f} \mid \mathbf{X}) = \mathcal{N}(\mathbf{f} \mid \boldsymbol{\mu}, \mathbf{K})$$

where $K_{ij} = K(\mathbf{x}_i, \mathbf{x}_j)$ and $\boldsymbol{\mu} = (m(\mathbf{x}_1), \dots, m(\mathbf{x}_N))$.

Regression using Gaussians

A Gaussian Process (GP) is denoted by:

$$f(\mathbf{x}) \sim GP(m(\mathbf{x}), \kappa(\mathbf{x}, \mathbf{x}'))$$

where $m(\mathbf{x})$ is the **mean function** and $K(\mathbf{x}, \mathbf{x}')$ is the **kernel** or **covariance function**, i.e.,

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$$

$$\kappa(\mathbf{x}, \mathbf{x}') = \mathbb{E} \left[(f(\mathbf{x}) - m(\mathbf{x})) (f(\mathbf{x}') - m(\mathbf{x}'))^T \right]$$

Gaussian processes in machine learning

- Gaussian processes (GPs), also known as Gaussian random fields.
- Originating in geostatistics, they were introduced by George Matheron around 1960 under the name kriging.
- Their use in machine learning has grown substantially since the 1990s.
- A Gaussian process extends the multivariate Gaussian distribution to an infinite collection of random variables indexed by input points.
- Its most common use in machine learning is probabilistic non-linear regression.
- Gaussian processes are also applied in pattern classification, dimensionality reduction, missing-data, multi-task learning, and Bayesian optimization.

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$f = \{f(x) : x \in \mathcal{X}\}.$$

Usually $f(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$

Gaussian processes

Understanding y requires only finite-dimensional distributions (FDDs):

- For any x_1, \dots, x_n and $n \in \mathbb{N}$:
 $\mathbb{P}(y(x_1) \leq c_1, \dots, y(x_n) \leq c_n).$
- These FDDs completely determine the law of y .

Gaussian processes

Understanding y requires only finite-dimensional distributions (FDDs):

- For any x_1, \dots, x_n and $n \in \mathbb{N}$:
 $\mathbb{P}(y(x_1) \leq c_1, \dots, y(x_n) \leq c_n).$
- These FDDs completely determine the law of y .

Gaussian process definition

A Gaussian process has Gaussian FDDs

$$(y(x_1), \dots, y(x_n)) \sim \mathcal{N}_n(\mu, \Sigma)$$

We write $y(\cdot) \sim \mathcal{GP}$ when y is a Gaussian process.

Mean and covariance function

To fully specify the law of a GP, we need to specify mean and covariance functions:

$$f(\mathbf{x}_1), \dots, f(\mathbf{x}_N) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')), \quad \forall \mathbf{x}, \mathbf{x}' \in \mathcal{X} \subset \mathbb{R}^d$$

where

[mean function]

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$$

[kernel covariance function]

$$\begin{aligned} k(\mathbf{x}, \mathbf{x}') &= \text{Cov}[f(\mathbf{x}), f(\mathbf{x}')] \\ &= \mathbb{E}[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))] \end{aligned}$$

Mean function $m(\mathbf{x})$

We can use any mean function we want $m(x) = \mathbb{E}[f(\mathbf{x})]$.

- $m(\mathbf{x}) = 0$
- $m(\mathbf{x}) = \text{const}$
- $m(x) = \phi(\mathbf{x})^\top \mathbf{w}$
- Neural networks.

Popular choices are:

Mean function $m(\mathbf{x})$

We can use any mean function we want $m(x) = \mathbb{E}[f(\mathbf{x})]$.

- $m(\mathbf{x}) = 0$
- $m(\mathbf{x}) = \text{const}$
- $m(x) = \phi(\mathbf{x})^\top \mathbf{w}$
- Neural networks.

Popular choices are:

Let $f(\mathbf{x}) = \phi(\mathbf{x})^\top \mathbf{w}$ with $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I})$.

Mean function $m(\mathbf{x})$

We can use any mean function we want $m(x) = \mathbb{E}[f(\mathbf{x})]$.

- $m(\mathbf{x}) = 0$
- $m(\mathbf{x}) = \text{const}$
- $m(x) = \phi(\mathbf{x})^\top \mathbf{w}$
- Neural networks.

Popular choices are:

Let $f(\mathbf{x}) = \phi(\mathbf{x})^\top \mathbf{w}$ with $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I})$.

Then $m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$

Mean function $m(\mathbf{x})$

We can use any mean function we want $m(x) = \mathbb{E}[f(\mathbf{x})]$.

- $m(\mathbf{x}) = 0$
- $m(\mathbf{x}) = \text{const}$
- $m(x) = \phi(\mathbf{x})^\top \mathbf{w}$
- Neural networks.

Popular choices are:

Let $f(\mathbf{x}) = \phi(\mathbf{x})^\top \mathbf{w}$ with $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I})$.

Then $m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})] = \mathbb{E}[\mathbf{w}^\top \phi(\mathbf{x})]$

Mean function $m(\mathbf{x})$

We can use any mean function we want $m(x) = \mathbb{E}[f(\mathbf{x})]$.

- $m(\mathbf{x}) = 0$
- $m(\mathbf{x}) = \text{const}$
- $m(x) = \phi(\mathbf{x})^\top \mathbf{w}$
- Neural networks.

Popular choices are:

Let $f(\mathbf{x}) = \phi(\mathbf{x})^\top \mathbf{w}$ with $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I})$.

Then $m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})] = \mathbb{E}[\mathbf{w}^\top \phi(\mathbf{x})] = \mathbb{E}[\mathbf{w}^\top] \mathbb{E}[\phi(\mathbf{x})]$

Mean function $m(\mathbf{x})$

We can use any mean function we want $m(x) = \mathbb{E}[f(\mathbf{x})]$.

- $m(\mathbf{x}) = 0$
- $m(\mathbf{x}) = \text{const}$
- $m(x) = \phi(\mathbf{x})^\top \mathbf{w}$
- Neural networks.

Popular choices are:

Let $f(\mathbf{x}) = \phi(\mathbf{x})^\top \mathbf{w}$ with $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I})$.

Then $m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})] = \mathbb{E}[\mathbf{w}^\top \phi(\mathbf{x})] = \mathbb{E}[\mathbf{w}^\top] \mathbb{E}[\phi(\mathbf{x})] = \mathbf{0}$

Kernel covariance function $k(\mathbf{x}, \mathbf{x}')$

The covariance function determines the nature of the GP, i.e., the hypothesis space/space of functions.

We usually use a covariance function that is a function of the indexes/locations:

$$k(\mathbf{x}, \mathbf{x}') = \text{Cov}(f(\mathbf{x}), f(\mathbf{x}')) ,$$

$k(\cdot, \cdot)$ must satisfy:

- Symmetry, i.e., $k(\mathbf{x}, \mathbf{x}') = k(\mathbf{x}', \mathbf{x})$.
- For any locations $\mathbf{x}_1, \dots, \mathbf{x}_N$, the $N \times N$ Gram matrix \mathbf{K} with $\mathbf{K}_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$ must be a positive semi-definite matrix.

Kernel covariance function $k(\mathbf{x}, \mathbf{x}')$

Let $f(\mathbf{x}) = \phi(\mathbf{x})^\top \mathbf{w}$ with $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \alpha^{-1} \mathbf{I})$.

The kernel covariance $k(\mathbf{x}, \mathbf{x}')$ is given by:

$$\begin{aligned} k(\mathbf{x}, \mathbf{x}') &= \mathbb{E}[f(\mathbf{x})f(\mathbf{x}')^\top] \\ &= \phi(\mathbf{x})^\top \mathbb{E}[\mathbf{w}\mathbf{w}^\top] \phi(\mathbf{x}') \\ &= \phi(\mathbf{x})^\top \frac{\mathbf{I}}{\alpha} \phi(\mathbf{x}') \\ &= \frac{\phi(\mathbf{x})^\top \phi(\mathbf{x}')}{\alpha} \end{aligned}$$

Kernel covariance function $k(\mathbf{x}, \mathbf{x}')$

A widely used kernel in GPs is the squared-exponential or RBF (radial basis function) kernel is given by:

$$k(x, x') = \sigma_f^2 \exp\left(-\frac{(x - x')^2}{2\ell^2}\right), \quad x \in \mathbb{R}$$

where σ_f^2 is the signal variance parameter and ℓ the length-scale parameter.

How do we draw samples from a GP?

- Given the mean function and covariance function for a GP, we can draw samples using a multivariate Gaussian distribution.
- To sample from the multivariate Gaussian distribution, we need a mean vector and a covariance matrix.
- The mean vector is obtained from the mean function.
- The covariance matrix is obtained from the covariance function.

Sampling from a GP, example

We have the set of x values:

$$\{x_1, x_2, \dots, x_N\} \subseteq \mathbb{R}$$

These are the indexes of the stochastic process.

We now compute the covariance matrix

$$\mathbf{K} = \begin{bmatrix} k(x_1, x_1) & k(x_1, x_2) & \cdots & k(x_1, x_N) \\ k(x_2, x_1) & k(x_2, x_2) & \cdots & k(x_2, x_N) \\ \vdots & \vdots & \vdots & \vdots \\ k(x_N, x_1) & k(x_N, x_2) & \cdots & k(x_N, x_N) \end{bmatrix}$$

We assume the mean function is constant and equal to zero, i.e.,
 $m(x) = 0 \forall x$.

Sampling from a GP, example

To generate functions from this GP, we will then sample from:


$$f(x_1), \dots, f(x_N)$$

$$\sim \mathcal{N} \left(\begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \begin{bmatrix} k(x_1, x_1) & k(x_1, x_2) & \cdots & k(x_1, x_N) \\ k(x_2, x_1) & k(x_2, x_2) & \cdots & k(x_2, x_N) \\ \vdots & \vdots & \vdots & \vdots \\ k(x_N, x_1) & k(x_N, x_2) & \cdots & k(x_N, x_N) \end{bmatrix} \right)$$


What we plot is x_i and $f(x_i)$, $\forall i$.



Sampling from a GP, example




$$\ell = 0.1, \sigma_f = 1$$




$\ell = 0.1, \sigma_f = 1$









$\ell = 0.1, \sigma_f = 1$




$\ell = 1, \sigma_f = 1$














$\ell = 3, \sigma_f = 1$


















$\ell = 0.1, \sigma_f = 2$

























$\ell = 0.1, \sigma_f = 1$  $\ell = 1, \sigma_f = 1$  $\ell = 3, \sigma_f = 1$  $\ell = 0.1, \sigma_f = 2$  $\ell = 1, \sigma_f = 2$  $\ell = 3, \sigma_f = 2$  $\ell = 0.1, \sigma_f = 3$ 


$\ell = 0.1, \sigma_f = 1$  $\ell = 1, \sigma_f = 1$  $\ell = 3, \sigma_f = 1$  $\ell = 0.1, \sigma_f = 2$  $\ell = 1, \sigma_f = 2$  $\ell = 3, \sigma_f = 2$  $\ell = 0.1, \sigma_f = 3$  $\ell = 1, \sigma_f = 3$ 

$\ell = 0.1, \sigma_f = 1$  $\ell = 1, \sigma_f = 1$  $\ell = 3, \sigma_f = 1$  $\ell = 0.1, \sigma_f = 2$  $\ell = 1, \sigma_f = 2$  $\ell = 3, \sigma_f = 2$  $\ell = 0.1, \sigma_f = 3$  $\ell = 1, \sigma_f = 3$  $\ell = 3, \sigma_f = 3$ 

$\ell = 0.1, \sigma_f = 1$  $\ell = 1, \sigma_f = 1$  $\ell = 3, \sigma_f = 1$  $\ell = 0.1, \sigma_f = 2$  $\ell = 1, \sigma_f = 2$  $\ell = 3, \sigma_f = 2$  $\ell = 0.1, \sigma_f = 3$  $\ell = 1, \sigma_f = 3$  $\ell = 3, \sigma_f = 3$ 


From prior to posterior

$$\ell = 1, \sigma_f = 1$$




From prior to posterior

$$\ell = 1, \sigma_f = 1$$



From prior to posterior

$$\ell = 1, \sigma_f = 1$$



From prior to posterior

Training set:

$$\mathcal{D} = \{(\mathbf{x}_i, f_i), i = 1 : N\}$$

Test set:

$$\{\mathbf{x}_i^*, i = 1 : N_*\}$$

The training and test sets are organized as follows:

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{pmatrix}$$

training inputs

We want to predict the function outputs \mathbf{f}^* .

From prior to posterior

Training set:

$$\mathcal{D} = \{(\mathbf{x}_i, f_i), i = 1 : N\}$$

Test set:

$$\{\mathbf{x}_i^*, i = 1 : N_*\}$$

The training and test sets are organized as follows:

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{pmatrix} \quad \mathbf{f} = \begin{pmatrix} f_1 \\ \vdots \\ f_N \end{pmatrix}$$

training inputs **training targets**

We want to predict the function outputs \mathbf{f}^* .

From prior to posterior

Training set:

$$\mathcal{D} = \{(\mathbf{x}_i, f_i), i = 1 : N\}$$

Test set:

$$\{\mathbf{x}_i^*, i = 1 : N_*\}$$

The training and test sets are organized as follows:

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{pmatrix} \quad \mathbf{f} = \begin{pmatrix} f_1 \\ \vdots \\ f_N \end{pmatrix} \quad \mathbf{X}^* = \begin{pmatrix} \mathbf{x}_1^{*\top} \\ \vdots \\ \mathbf{x}_{N_*}^{*\top} \end{pmatrix}$$

training inputs training targets test inputs

We want to predict the function outputs \mathbf{f}^* .

From prior to posterior

Training set:

$$\mathcal{D} = \{(\mathbf{x}_i, f_i), i = 1 : N\}$$

Test set:

$$\{\mathbf{x}_i^*, i = 1 : N_*\}$$

The training and test sets are organized as follows:

$$\mathbf{X} = \begin{pmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_N^\top \end{pmatrix}$$

$$\mathbf{f} = \begin{pmatrix} f_1 \\ \vdots \\ f_N \end{pmatrix}$$

$$\mathbf{X}^* = \begin{pmatrix} \mathbf{x}_1^{*\top} \\ \vdots \\ \mathbf{x}_{N_*}^{*\top} \end{pmatrix}$$

$$\mathbf{f}^* = \begin{pmatrix} f_1^* \\ \vdots \\ f_{N_*}^* \end{pmatrix}$$

training inputs

training targets

test inputs

predictions

We want to predict the function outputs \mathbf{f}^* .

From prior to posterior

The prior GP joint distribution of \mathbf{f} and \mathbf{f}^* is multivariate normal, i.e.:

$$\begin{pmatrix} \mathbf{f} \\ \mathbf{f}_* \end{pmatrix} \sim \mathcal{N} \left(\begin{pmatrix} \boldsymbol{\mu} \\ \boldsymbol{\mu}_* \end{pmatrix}, \begin{pmatrix} \mathbf{K} & \mathbf{K}_* \\ \mathbf{K}_*^\top & \mathbf{K}_{**} \end{pmatrix} \right)$$

where:

\mathbf{K} is of size $N \times N$, with $[\mathbf{K}]_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$.

\mathbf{K}_* is of size $N \times N_*$, with $[\mathbf{K}_*]_{ij} = k(\mathbf{x}_i, \mathbf{x}_{*j})$.

\mathbf{K}_{**} is of size $N_* \times N_*$, with $[\mathbf{K}_{**}]_{ij} = k(\mathbf{x}_{*i}, \mathbf{x}_{*j})$.

From prior to posterior


Conditioning on the training targets \mathbf{f} gives the posterior GP,
i.e.:

$$\mathbf{f}_* \mid \mathbf{X}_*, \mathbf{X}, \mathbf{f} \sim \mathcal{N}(\boldsymbol{\mu}_*, \boldsymbol{\Sigma}_*)$$


The posterior has the following form:

$$\boldsymbol{\mu}_* = \boldsymbol{\mu}(\mathbf{X}_*) + \mathbf{K}_*^\top \mathbf{K}^{-1} (\mathbf{f} - \boldsymbol{\mu})$$

$$\boldsymbol{\Sigma}_* = \mathbf{K}_{**} - \mathbf{K}_*^\top \mathbf{K}^{-1} \mathbf{K}_*$$




(a)




(b)

- (a) Samples from the prior, $p(\mathbf{f} \mid \mathbf{X})$, using a squared exponential kernel.
- (b) Samples from a GP posterior, $p(\mathbf{f}_* \mid \mathbf{X}_*, \mathbf{X}, \mathbf{f})$.


Numerical example



Numerical example



Numerical example



Training samples:

$$(x_1, f(x_1)) = (-2, -0.62)$$
$$(x_2, f(x_2)) = (1, 0.54)$$

$$x^* = 0.5$$

$$f^* = ?$$

Numerical example

Suppose f is a Gaussian process, then

$$\underbrace{f(x_1), f(x_2)}_{\text{training samples}}, \underbrace{f(x^*)}_{\text{test sample}} \sim \mathcal{N}(\mathbf{0}, \Sigma)$$

where:

$$\Sigma = \left(\begin{array}{cc|c} k(x_1, x_1) & k(x_1, x_2) & k(x_1, x^*) \\ k(x_2, x_1) & k(x_2, x_2) & k(x_2, x^*) \\ \hline k(x^*, x_1) & k(x^*, x_2) & k(x^*, x^*) \end{array} \right) = \left(\begin{array}{cc|c} 1 & 0.01 & 0.04 \\ 0.01 & 1 & 0.88 \\ \hline 0.04 & 0.88 & 1 \end{array} \right)$$

The value of the function f^* at the testing point $x^* = 0.5$ is computed from $p(\mathbf{f}_* | \mathbf{f}) = \mathcal{N}(\boldsymbol{\mu}_*, \Sigma_*)$ with:

$$\begin{aligned} \boldsymbol{\mu}_* &= \mathbf{K}_*^\top \mathbf{K}^{-1} (\mathbf{f} - \boldsymbol{\mu}) \\ &= \begin{pmatrix} k(x^*, x_1) & k(x^*, x_2) \end{pmatrix} \begin{pmatrix} k(x_1, x_1) & k(x_1, x_2) \\ k(x_2, x_1) & k(x_2, x_2) \end{pmatrix}^{-1} \begin{pmatrix} f(x_1) \\ f(x_2) \end{pmatrix} \\ &= 0.46 \end{aligned}$$

Numerical example

Suppose f is a Gaussian process, then

$$\underbrace{f(x_1), f(x_2)}_{\text{training samples}}, \underbrace{f(x^*)}_{\text{test sample}} \sim \mathcal{N}(0, \Sigma)$$


where:

$$\Sigma = \left(\begin{array}{cc|c} k(x_1, x_1) & k(x_1, x_2) & k(x_1, x^*) \\ k(x_2, x_1) & k(x_2, x_2) & k(x_2, x^*) \\ \hline k(x^*, x_1) & k(x^*, x_2) & k(x^*, x^*) \end{array} \right) = \left(\begin{array}{cc|c} 1 & 0.01 & 0.04 \\ 0.01 & 1 & 0.88 \\ \hline 0.04 & 0.88 & 1 \end{array} \right)$$


The value of the function f^* at the testing point $x^* = 0.5$ is computed from $p(f_* | f) = \mathcal{N}(\mu_*, \sigma^2_*)$ with:

$$\begin{aligned} \Sigma_* &= K_{**} - K_*^T K^{-1} K_* \\ &= k(x^*, x^*) - \begin{pmatrix} k(x^*, x_1) & k(x^*, x_2) \end{pmatrix} \begin{pmatrix} k(x_1, x_1) & k(x_1, x_2) \\ k(x_2, x_1) & k(x_2, x_2) \end{pmatrix}^{-1} \begin{pmatrix} k(x^*, x_1) \\ k(x^*, x_2) \end{pmatrix} \\ &= 0.22 \end{aligned}$$


Numerical example




Numerical example



Numerical example



Numerical example



Predictions using noisy observations

Noisy version of the underlying function:

$$y = f(\mathbf{x}) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma_y^2)$$

The covariance of the observed noisy responses is:

$$\text{cov}[y_p, y_q] = K(\mathbf{x}_p, \mathbf{x}_q) + \sigma_y^2 \mathbb{I}(p = q)$$

$$\text{cov}[\mathbf{y} \mid \mathbf{X}] = \mathbf{K} + \sigma_y^2 \mathbf{I}_N \triangleq \mathbf{K}_y$$

Predictions using noisy observations

The joint density of training and test samples is given by:

$$y_1, y_2, \dots, y_N, f_* \sim \mathcal{N} \left(\mathbf{0}, \begin{pmatrix} \mathbf{K}_y & \mathbf{k}_* \\ \mathbf{k}_*^\top & k_{**} \end{pmatrix} \right)$$

$$\begin{pmatrix} \mathbf{y} \\ \mathbf{f}_* \end{pmatrix} \sim \mathcal{N} \left(\mathbf{0}, \begin{pmatrix} \mathbf{K}_y & \mathbf{K}_* \\ \mathbf{K}_*^\top & \mathbf{K}_{**} \end{pmatrix} \right)$$

Predictions using noisy observations

The posterior predictive density is:

$$\begin{aligned} p(f_* \mid X_*, X, y) &= \mathcal{N}(f_* \mid \mu_*, \Sigma_*) \\ \mu_* &= K_*^T K_y^{-1} y \\ \Sigma_* &= K_{**} - K_*^T K_y^{-1} K_* \end{aligned}$$


In the case of a single test input, this simplifies as follows:

$$p(f_* \mid \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \mathcal{N}(f_* \mid \mathbf{k}_*^T \mathbf{K}_y^{-1} \mathbf{y}, k_{**} - \mathbf{k}_*^T \mathbf{K}_y^{-1} \mathbf{k}_*)$$


where $\mathbf{k}_* = [\kappa(\mathbf{x}_*, \mathbf{x}_1), \dots, \kappa(\mathbf{x}_*, \mathbf{x}_N)]$ and $k_{**} = \kappa(\mathbf{x}_*, \mathbf{x}_*)$. Another way to write the posterior mean is as follows:

$$\bar{f}_* = \mathbf{k}_*^T \mathbf{K}_y^{-1} \mathbf{y} = \sum_{i=1}^N \alpha_i \kappa(\mathbf{x}_i, \mathbf{x}_*)$$


where $\alpha = \mathbf{K}_y^{-1} \mathbf{y}$.



(a)



(b)



(c)

Which kernel functions are appropriate for a given problem?

How does the choice of different hyperparameter values affect the resulting model?

Kernels in Gaussian Processes

A **kernel** (or covariance function) $k(\mathbf{x}, \mathbf{x}')$ is the foundation of a Gaussian process. It defines the covariance structure between function values at any two input points, completely specifying the GP's prior distribution and determining:

- **Smoothness** of sample functions
- **Length-scale** (characteristic distance over which correlations decay)
- **Amplitude** (overall vertical scale of variations)
- **Periodicity** (if applicable)
- **Non-stationarity** (if applicable)

Kernels must be **positive semi-definite** to ensure valid covariance matrices.

Kernel Effects

1. Signal Variance (σ_f^2)

- Controls the **amplitude** or overall vertical scale of function variations.
- Larger values: function can vary over a wider range.
- Smaller values: function values stay closer to the mean.

2. Length-Scale (ℓ)

- Controls how quickly correlations **decay with distance**.
- **Large ℓ :** Points remain strongly correlated over longer distances.
Sample functions are smooth and slowly varying.
- **Small ℓ :** Correlations decay rapidly, function can vary more rapidly and rougher behavior.

Two common kernels

RBF (squared exponential) kernel:


$$k_{\text{RBF}}(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right).$$

Matérn kernel (general $\nu > 0$) can be written as:


$$k_{\text{Matérn}}(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \frac{1}{\Gamma(\nu) 2^{\nu-1}} \left(\frac{\sqrt{2\nu}}{\ell} \|\mathbf{x} - \mathbf{x}'\| \right)^\nu K_\nu \left(\frac{\sqrt{2\nu}}{\ell} \|\mathbf{x} - \mathbf{x}'\| \right),$$

where K_ν is the modified Bessel function of the second kind, ν controls smoothness, ℓ is the length scale, and σ_f^2 is the signal variance.


RBF kernel - samples from prior




RBF kernel - samples from prior




RBF kernel - samples from prior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - samples from posterior




RBF kernel - prior / posterior



Rational-quadratic kernel - prior / posterior



Matern kernel - prior / posterior



3. Stationarity

- **Stationary kernel:** Covariance depends only on the *difference* $x - x'$, invariant under translation
- **Non-stationary kernel:** Covariance depends on absolute locations, allowing position-dependent behavior

4. Isotropy

- **Isotropic kernel:** Covariance depends only on the *Euclidean distance* $\|x - x'\|$
- **Anisotropic kernel:** Different characteristic length-scales and couplings across dimensions

Stationary Kernels

A stationary covariance function satisfies:

$$\text{Cov}[f(\mathbf{x}), f(\mathbf{x}')] = k(\mathbf{x} - \mathbf{x}') = \psi(\mathbf{x} - \mathbf{x}')$$

Translation invariant: $k(\mathbf{x} + \mathbf{a}, \mathbf{x}' + \mathbf{a}) = k(\mathbf{x}, \mathbf{x}')$ for all shifts \mathbf{a} .

Example: RBF kernel

Isotropic Kernels

An isotropic kernel depends only on distance:

$$k(\mathbf{x}, \mathbf{x}') = \phi(\|\mathbf{x} - \mathbf{x}'\|)$$

where $\phi : [0, \infty) \rightarrow \mathbb{R}$ is a scalar function.

Property: Every isotropic kernel is stationary, but not all stationary kernels are isotropic.

Example: RBF kernel is isotropic when $\Lambda = \ell^2 I$ (identity matrix)

Anisotropic Kernels

A stationary *but not isotropic* kernel depends on the full lag vector $\mathbf{h} = \mathbf{x} - \mathbf{x}'$, including directional information:

$$k(\mathbf{x}, \mathbf{x}') = \sigma^2 \exp\left(-\frac{1}{2}(\mathbf{x} - \mathbf{x}')^\top \Lambda^{-1} (\mathbf{x} - \mathbf{x}')\right)$$

where Λ is a positive definite matrix defining different characteristic length-scales and couplings across dimensions.

Use case: When different input dimensions have different “importance” or correlation decay rates.

Non-Stationary Kernels

Covariance depends on absolute locations, not just their difference. Allow position-dependent behavior and varying smoothness.

Examples: Polynomial kernel, neural network kernel

Kernel: Squared-Exponential (RBF)

1D Form:

$$k(x, x') = \sigma_f^2 \exp\left(-\frac{(x - x')^2}{2\ell^2}\right)$$

Multi-dimensional isotropic form:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\ell^2}\right)$$

Characteristics: Infinitely differentiable (very smooth sample functions), **stationary and isotropic**, rapid decay with distance, widely used due to simplicity and smoothness.

Parameters: σ_f^2 , signal variance (amplitude); ℓ , length-scale.

Kernel: Matérn

General form:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\sqrt{2\nu} \frac{\|\mathbf{x} - \mathbf{x}'\|}{\ell} \right)^\nu K_\nu \left(\sqrt{2\nu} \frac{\|\mathbf{x} - \mathbf{x}'\|}{\ell} \right)$$

where K_ν is a modified Bessel function, $\nu > 0$ controls smoothness, and ℓ is the length-scale.

Common versions: $\nu = 1/2$, exponential kernel; $\nu = 3/2$; $\nu \rightarrow \infty$, approaches RBF kernel.

Characteristics: More flexible smoothness control than RBF, includes exponential kernel as special case, stationary and isotropic.

Parameters: σ_f^2 , signal variance; ℓ , length-scale; ν , smoothness parameter.

Kernel: Rational Quadratic (RQ)

Form:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \left(1 + \frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\alpha\ell^2} \right)^{-\alpha}$$

Characteristics: Mixture of RBF kernels with different length-scales; can be viewed as weighted superposition of exponential kernels at different scales, **stationary and isotropic**.

Parameters: σ_f^2 , signal variance, ℓ , length-scale; α , mixing parameter (larger α , closer to RBF).

Kernel: Periodic

Form:

$$k(\mathbf{x}, \mathbf{x}') = \sigma_f^2 \exp\left(-\frac{2 \sin^2(\pi |\mathbf{x} - \mathbf{x}'|/p)}{\ell^2}\right)$$

Characteristics: Captures periodic patterns in data; essential for time-series with seasonal patterns.

Parameters: σ_f^2 : signal variance; ℓ : length-scale; p : period (wavelength).

Kernel: Polynomial

Form:

$$k(\mathbf{x}, \mathbf{x}') = (\sigma_f^2 \mathbf{x} \cdot \mathbf{x}' + c)^d$$

or

$$k(\mathbf{x}, \mathbf{x}') = (\sigma_f^2 \mathbf{x} \cdot \mathbf{x}' + c)^d + \sigma_n^2$$

Characteristics: Often used with small d (typically 1, 2, or 3),
non-stationary kernel.

Parameters: σ_f^2 , signal variance, c , offset, d , degree (model complexity), σ_n^2 , noise variance (if included).

Kernel: White Noise

Form:

$$k(x, x') = \sigma_n^2 \delta(x - x')$$

where $\delta(\cdot)$ is the Dirac delta function.

In practice (discrete inputs with finite precision):

$$k_{\text{white noise}}(x, x') = \begin{cases} \sigma_n^2 & \text{if } x = x' \\ 0 & \text{otherwise} \end{cases}$$

Kernel: White Noise

- **Pure noise:** No correlation between any two distinct points.
- **Variance only on diagonal:** Contributes σ_n^2 to the diagonal of the covariance matrix.
- **Numerical stability:** Often added to smooth kernels to stabilize matrix inversion.
- **Independent observations:** Assumes each observation contains independent measurement noise.

Kernel: White Noise

1. **Likelihood noise term:** Observations $y_i = f(\mathbf{x}_i) + \epsilon_i$, where $\epsilon_i \sim \mathcal{N}(0, \sigma_n^2)$
2. **Jitter for numerical stability:** Added to other kernels to improve matrix conditioning:

$$k_{\text{total}} = k_{\text{smooth}}(\mathbf{x}, \mathbf{x}') + \sigma_n^2 \delta(\mathbf{x} - \mathbf{x}')$$

3. **Model white noise in data:** When you expect uncorrelated noise in observations

Kernel combinations

Kernels are closed under addition and multiplication:

- **Sum (additive):** $k_{\text{total}} = k_1 + k_2$ (combines features)
 - Example: RBF + Periodic (smooth + periodic patterns)
 - Example: RBF + White Noise (smooth signal + independent noise)
- **Product (multiplicative):** $k_{\text{total}} = k_1 \times k_2$ (modulates one kernel by another)
 - Example: (RBF) \times (Periodic) (smooth periodic trends)

Estimating the kernel parameters

$$K_y(x_p, x_q) = \sigma_f^2 \exp\left(-\frac{1}{2\ell^2} (x_p - x_q)^2\right) + \sigma_y^2 \delta_{pq}$$

We can maximize the marginal likelihood:

$$p(\mathbf{y} \mid \mathbf{X}) = \int p(\mathbf{y} \mid \mathbf{f}, \mathbf{X}) p(\mathbf{f} \mid \mathbf{X}) d\mathbf{f}$$

Estimating the kernel parameters

We can maximize the marginal likelihood:

$$p(\mathbf{y} \mid \mathbf{X}) = \int p(\mathbf{y} \mid \mathbf{f}, \mathbf{X}) p(\mathbf{f} \mid \mathbf{X}) d\mathbf{f}$$

where:

$$p(\mathbf{f} \mid \mathbf{X}) = \mathcal{N}(\mathbf{f} \mid \mathbf{0}, \mathbf{K})$$

$$p(\mathbf{y} \mid \mathbf{f}) = \prod_i \mathcal{N}(y_i \mid f_i, \sigma_y^2)$$

Estimating the kernel parameters

The marginal likelihood is given by:

$$\begin{aligned}\log p(\mathbf{y} \mid \mathbf{X}) &= \log \mathcal{N}(\mathbf{y} \mid \mathbf{0}, \mathbf{K}_y) \\ &= -\underbrace{\frac{1}{2}\mathbf{y}\mathbf{K}_y^{-1}\mathbf{y}}_{\text{data fit}} - \underbrace{\frac{1}{2}\log|\mathbf{K}_y|}_{\text{model complexity}} - \underbrace{\frac{N}{2}\log(2\pi)}_{\text{constant}}\end{aligned}$$

Estimating the kernel parameters


Let the kernel parameters (also called hyper-parameters) be denoted by θ .

One can show that:

$$\begin{aligned}\frac{\partial}{\partial \theta_j} \log p(\mathbf{y} \mid \mathbf{X}) &= \frac{1}{2} \mathbf{y}^T \mathbf{K}_y^{-1} \frac{\partial \mathbf{K}_y}{\partial \theta_j} \mathbf{K}_y^{-1} \mathbf{y} - \frac{1}{2} \text{tr} \left(\mathbf{K}_y^{-1} \frac{\partial \mathbf{K}_y}{\partial \theta_j} \right) \\ &= \frac{1}{2} \text{tr} \left((\boldsymbol{\alpha} \boldsymbol{\alpha}^T - \mathbf{K}_y^{-1}) \frac{\partial \mathbf{K}_y}{\partial \theta_j} \right)\end{aligned}$$

where $\boldsymbol{\alpha} = \mathbf{K}_y^{-1} \mathbf{y}$.

Log marginal likelihood



Computational cost

- One difficulty with GPs is the computational cost of training them: $O(n^3)$ (and $O(n^2)$ memory).
- They work our of the box for n in the order of a few thousands.
- There are many ways to side-step this cost: inducing inputs, efficient matrix-vector multiplications, random features, etc.
- These days we can use GPs for n in the order of tens of millions.